

Linear Model

Libraries

```
In [1]: import pandas as pd
import numpy as np
import os
from functools import reduce
from datetime import date, datetime

from scipy.stats import ks_2samp

from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_absolute_error
from sklearn.metrics import roc_auc_score, accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

import matplotlib.pyplot as plt
import seaborn as sns
import pygal
from scikitplot.metrics import plot_roc_curve

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
pd.set_option('display.max_columns', 500)
```

Data Reading

Select specific columns to save memory and processing power

```
In [2]: csv = '2018_Yellow_Taxi_Trip_Data.csv'

fields = ['VendorID', 'tpep_pickup_datetime',
          'passenger_count', 'trip_distance',
          'PULocationID', 'payment_type',
          'total_amount']

models = []

In [3]: iterator = pd.read_csv(csv, usecols=fields,
                               dtype={'VendorID': 'int8', 'passenger_count': 'int8',
                                      'trip_distance': 'float32', 'PULocationID': 'int16',
                                      'payment_type': 'int8', 'total_amount': 'float32'},
                               parse_dates=["tpep_pickup_datetime"],
                               chunksize=1000000)
```

Functions

Data Filter

```
In [4]: def data_filter(df):
''' data_filter

Filters data where:
-Data is not from 2018
-Total amount is less than zero
-Passenger count is less than one
-Trip distance is less or equal to zero

-Drop passenger_count (we dont use it later in the model)

Parameter
-----
df: Dataframe chunk
'''

#Filters data not from 2018
start_date = pd.to_datetime('2018-01-01 00:00:00')
end_date = pd.to_datetime('2019-01-01 00:00:00')
mask = (df['tpep_pickup_datetime'] >= start_date) & (df['tpep_pickup_datetime'] < end_date)
df = df.loc[mask]

#Filters data where total amount is less than zero
df = df[df["total_amount"] > 0.00]

#Filters where passenger count is less than one
df = df[df["passenger_count"] >= 1]

#Filters where trip distance = 0
df = df[df["trip_distance"] > 0]

df.drop(['passenger_count'], axis=1, inplace = True)

return df
```

```
In [5]: def dates(df):
'''dates

Splits Datetime into two new columns:
-pickup_date
-pickup_hour

Adds column:
-pickup_hour where: First two digits represent month,
                    Middle two represent day of the week
                    Last two represent the pickup hour

Drops the original datetime column

Parameter
-----
df: Dataframe chunk
'''

df['pickup_date'] = df['tpep_pickup_datetime'].dt.date

df['pickup_time'] = df['tpep_pickup_datetime'].dt.time
df['pickup_time'] = df['pickup_time'].astype(str)
```

```

##df['month'] = pd.to_datetime(df['tpep_pickup_datetime'].dt.month)
df['month'] = df['tpep_pickup_datetime'].dt.strftime('%m')

#Day of week {01: Monday .... 07:Sunday}
df['pickup_day_of_week'] = df['tpep_pickup_datetime'].dt.dayofweek +1
df['pickup_day_of_week'] = '0' + df['pickup_day_of_week'].astype(str)

df['hour'] = df['tpep_pickup_datetime'].dt.strftime('%H')

df["pickup_hour"] = df['month'] + df['pickup_day_of_week'] + df['hour']

df.drop(["tpep_pickup_datetime",
        "month",
        "pickup_day_of_week",
        "hour"], axis=1, inplace=True)

return df

```

Hours Catalog

```

In [6]: def hours_catalog(df):
        '''hours_catalog

        Creates id_hour (int16) column based on pickup_hour sorted column

        Parameter
        -----
        df: Chunk of Dataframe that contains pickup_hour values:
            First two values: Month number
            Third and Fourth values: Day of the week {01: Monday .... 07:Sunday}
            Last two values: Hour

        '''
        catfh = df[['pickup_hour']].drop_duplicates().sort_values('pickup_hour', asce
        catfh["id_hour"] = (catfh.index+1)
        catfh["id_hour"] = catfh["id_hour"].astype('int16')

        df = df.merge(catfh, on='pickup_hour', how='inner')
        df.drop('pickup_hour', axis=1, inplace=True)

        return df

```

Data Engineering

```

In [7]: def trans(df, ancla, k):
        '''trans

        Parameters
        -----
        df: Chunk of Dataframe
        ancla: current ancla to run
        k: current step
        '''

        aux = df.loc[(df['id_hour']>=(ancla-k+1))&(df['id_hour']<=ancla)].reset_inde
        aux = aux[['id_hour', 'trip_distance', 'total_amount', 'pickup_time', 'PULoca
        aux['hora'] = aux['pickup_time'].map(lambda x:int(x.split(':')[0])//6).astyp
        aux.drop('pickup_time', axis=1, inplace=True)
        aux['n'] = 1

```

```

t = aux.copy()

for v in vard:
    t[v] = 'total_%s'%v
aux = pd.concat([aux,t],ignore_index=True)

def piv(aux,v,ancla):
    aux = aux.pivot_table(index='PULocationID',
                           columns=v,
                           values=varc,
                           aggfunc=['min','max','mean','sum','std'])

    aux.columns = ["v_%s"%v+"_".join(x)+"_%d"%k for x in aux.columns]
    return aux.reset_index().assign(ancla=ancla)

aux = reduce(lambda x,y:pd.merge(x,y,on=um,how='outer'),map(lambda v:piv(aux
return aux

```

TAD & Target

```

In [8]: def TAD(X, y, um):

        tad = X.merge(y,on=um,how='inner')

        return tad

```

```

In [9]: def target(df, ancla, vdes):
        aux = df.loc[(df['id_hour']>ancla)&(df['id_hour']<=(ancla+vdes))].reset_index()
        aux = aux.groupby(um[0]).sum()
        return aux.assign(ancla=ancla).reset_index()

```

Impute

```

In [10]: def impute(tad, varc):
        '''impute

        Returns
        X: Based on TAD column with only columns from varc list
        Xi: Imputed df based on X with median strategy

        Parameters
        .....
        tad: TAD table
        varc: list of continued variables

        ...

        X = tad[varc].copy()
        im = SimpleImputer(strategy='median')
        im.fit(X)
        Xtrans = im.transform(X)

        Xi = pd.DataFrame(Xtrans, columns=varc)

        #Kolmogorov-Smirnov (Two distributions are statistically equal)
        ks = pd.DataFrame(map(lambda v: (v,ks_2samp(tad[v].dropna(),Xi[v])).statistic
        print(ks.loc[ks['ks']>.1])

```

```
return X, Xi
```

Extreme Values

```
In [11]: def extremo(df,v,ci,cs):
    aux = df[um+[v]].copy()
    aux['ol_%s'%v] = ((aux[v]<ci)|(aux[v]>cs)).astype(int)
    return aux.drop(v,axis=1)

def extreme(X, um, tad):

    cotas = X.describe(percentiles=[0.01,0.99]).T[['1%', '99%']].reset_index().v
    ext = reduce(lambda x,y:pd.merge(x,y,on=um,how='outer'),map(lambda z:extremo
    varol = [v for v in ext if v[:2]=='ol']
    ext['extremo'] = ext[varol].max(axis=1)
    print(ext['extremo'].describe())
```

Modelling (Find parameters $\vec{\theta}$ of model f for

$$\vec{y} = f(\mathcal{X}) \rightarrow y = \theta_0 + \vec{\theta} \cdot \vec{x}$$

```
In [12]: def regression(tad, tgt, Xi):
    '''regression

    Creates a linear model for the df passed.
    Appends the model to the models list so we can dump it with pickle

    Parameters
    -----
    tad: TAD table for the dataframe chunk
    tgt: target column
    Xi: table with imputed values

    ...

    y = tad[tgt].copy()
    Xt,Xv,yt,yv = train_test_split(Xi,y,train_size=0.7)
    modelo = LinearRegression()
    hiperparametros=dict(fit_intercept=[True,False],normalize=[True,False])

    grid = GridSearchCV(param_grid=hiperparametros,
                        estimator=modelo,
                        cv=10,
                        scoring='neg_mean_absolute_error',
                        n_jobs=-1,
                        verbose=True)

    grid.fit(Xt,yt)

    modelo = grid.best_estimator_
    modelo.fit(Xt,yt)
    modelo.intercept_

    print(mean_absolute_error(y_true=yt,y_pred=modelo.predict(Xt)))
    print(mean_absolute_error(y_true=yv,y_pred=modelo.predict(Xv)))

    Xv['y^'] = modelo.predict(Xv)
```

```
Xv['y'] = yv
modelo.coef_

models.append(modelo)
```

First Chunk

Get the first chunk of data to test our functions

```
In [13]: %%time
df = iterator.get_chunk()
```

CPU times: user 1min 18s, sys: 26.2 ms, total: 1min 18s
Wall time: 1min 19s

```
In [14]: df.isnull().sum()
```

```
Out[14]: VendorID                0
tpep_pickup_datetime          0
passenger_count               0
trip_distance                 0
PULocationID                 0
payment_type                  0
total_amount                  0
dtype: int64
```

No missing data, no need to worry. We next filter the data for 2018.

```
In [15]: %%time
df = data_filter(df)
df = dates(df)
```

CPU times: user 11.6 s, sys: 318 ms, total: 11.9 s
Wall time: 11.9 s

Hours Catalog

```
In [16]: df = hours_catalog(df)
df.head()
```

```
Out[16]:
```

	VendorID	trip_distance	PULocationID	payment_type	total_amount	pickup_date	pickup_time	id_l
0	2	8.00	230	2	26.80	2018-09-22	23:46:37	
1	2	1.70	141	1	11.76	2018-09-22	23:00:43	
2	2	3.84	163	2	15.30	2018-09-22	23:10:32	
3	2	1.80	166	2	9.80	2018-09-22	23:27:25	
4	2	2.43	229	1	13.30	2018-09-22	23:50:39	

Data Engineering

```
In [17]: horai, horaf = df[['id_hour']].describe().T[['min', 'max']].values[0].tolist()
horai, horaf
```

```
Out[17]: (1.0, 134.0)
```

```
In [18]: vobs = 24
vdes = 1
anclai = int(horai)+vobs-1
anclaf = int(horaf)-vdes
anclai, anclaf
```

```
Out[18]: (24, 133)
```

```
In [19]: um = ['PULocationID', 'ancla']
ancla = 24
step = 4
varc = ['trip_distance', 'total_amount', 'n']
vard = ['hora']
```

```
In [20]: %%time
X_test = pd.concat(map(lambda ancla: reduce(lambda x, y: pd.merge(x, y, on=um, how='outer'),
map(lambda k: trans(df, ancla, k), range(step, vobs+step, step))), range(anclai,
CPU times: user 4min 58s, sys: 2.82 s, total: 5min 1s
Wall time: 5min 1s
```

```
In [21]: X_test.head()
```

```
Out[21]:
```

	PULocationID	v_hora_min_n_3_4	v_hora_min_n_total_hora_4	v_hora_min_total_amount_3_4	v_hora_min_total_hora_4
0	1	1.0	1.0	105.30	
1	4	1.0	1.0	5.80	
2	7	1.0	1.0	4.80	
3	10	1.0	1.0	35.38	
4	12	1.0	1.0	7.55	

```
In [22]: y_test= pd.concat(map(lambda ancla: target(df, ancla, vdes), range(anclai, anclaf+1)))
y_test.head()
```

```
Out[22]:
```

	PULocationID	y	ancla
0	4	7	24
1	7	21	24
2	10	6	24
3	12	1	24
4	13	27	24

TAD (Tabla analítica de datos) $\vec{y} = f(\mathcal{X})$

```
In [23]: tad_test = TAD(X_test, y_test, um)
tad_test.head()
```

```
Out[23]:
```

	PULocationID	v_hora_min_n_3_4	v_hora_min_n_total_hora_4	v_hora_min_total_amount_3_4	v_hora_min_total_hora_4
0	4	1.0	1.0	5.80	
1	7	1.0	1.0	4.80	

	PULocationID	v_hora_min_n_3_4	v_hora_min_n_total_hora_4	v_hora_min_total_amount_3_4	v_hora
2	10	1.0	1.0	35.38	
3	12	1.0	1.0	7.55	
4	13	1.0	1.0	0.31	

Exploratory Analysis

Variable Selection

```
In [24]: varc_test = [v for v in tad_test.columns if v[:2]=="v_"]
        tgt_test = 'y'
```

Missing Values Counting

```
In [46]: miss_test = 1-tad_test[varc_test].describe().T[['count']]/len(tad_test)
        miss_test.describe()
```

```
Out[46]:
```

	count
count	450.000000
mean	0.293301
std	0.224234
min	0.000000
25%	0.113302
50%	0.231716
75%	0.467293
max	0.736255

Impute

```
In [26]: X_test, Xi_test = impute(tad_test, varc_test)
```

	variable	ks
2	v_hora_min_total_amount_3_4	0.329459
4	v_hora_min_trip_distance_3_4	0.345802
8	v_hora_max_total_amount_3_4	0.345092
10	v_hora_max_trip_distance_3_4	0.346276
14	v_hora_mean_total_amount_3_4	0.346986
..
444	v_hora_sum_n_2_4	0.328851
445	v_hora_sum_total_amount_2_4	0.328851
446	v_hora_sum_trip_distance_2_4	0.328851
448	v_hora_std_total_amount_2_4	0.344966
449	v_hora_std_trip_distance_2_4	0.344966

[181 rows x 2 columns]

Valores Extremos

```
In [27]: extreme(X_test, um, tad_test)
```



```
count    9585.000000
mean      0.509442
std       0.499937
min       0.000000
25%       0.000000
50%       1.000000
75%       1.000000
max       1.000000
Name: extremo, dtype: float64
```

Modelación (encontrar los parámetros $\vec{\theta}$ del modelo f
para $\vec{y} = f(\mathcal{X}) \rightarrow y = \theta_0 + \vec{\theta} \cdot \vec{x}$

In [28]: `tad_test.head()`

```
Out[28]:
```

	PULocationID	v_hora_min_n_3_4	v_hora_min_n_total_hora_4	v_hora_min_total_amount_3_4	v_hora
0	4	1.0	1.0	5.80	
1	7	1.0	1.0	4.80	
2	10	1.0	1.0	35.38	
3	12	1.0	1.0	7.55	
4	13	1.0	1.0	0.31	

```
In [29]: y_test = tad_test['y'].copy()
Xt_test, Xv_test, yt_test, yv_test = train_test_split(Xi_test, y_test, train_size=0.7)
modelo_test = LinearRegression()
hiperparametros=dict(fit_intercept=[True, False], normalize=[True, False])

grid_test = GridSearchCV(param_grid=hiperparametros,
                          estimator=modelo_test,
                          cv=10,
                          scoring='neg_mean_absolute_error',
                          n_jobs=-1,
                          verbose=True)
```

In [30]: `grid_test.fit(Xt_test, yt_test)`

```
Fitting 10 folds for each of 4 candidates, totalling 40 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 40 out of 40 | elapsed: 7.7s finished
```

```
Out[30]: GridSearchCV(cv=10, estimator=LinearRegression(), n_jobs=-1,
                      param_grid={'fit_intercept': [True, False],
                                   'normalize': [True, False]},
                      scoring='neg_mean_absolute_error', verbose=True)
```

In [31]: `modelo_test = grid_test.best_estimator_`

In [32]: `modelo_test.fit(Xt_test, yt_test)`

```
Out[32]: LinearRegression(fit_intercept=False, normalize=True)
```

In [33]: `modelo_test.intercept_`

Out[33]: 0.0


```
In [34]: print(mean_absolute_error(y_true=yt_test,y_pred=modelo_test.predict(Xt_test)))
print(mean_absolute_error(y_true=yv_test,y_pred=modelo_test.predict(Xv_test)))
```

22.515774748259656

22.701732484535974


```
In [35]: sns.distplot(modelo_test.predict(Xt_test),hist=False,kde_kws={'cumulative':True})
sns.distplot(yt_test,hist=False,kde_kws={'cumulative':True})
```

Out[35]: <AxesSubplot:xlabel='y', ylabel='Density'>



```
In [36]: sns.distplot(modelo_test.predict(Xv_test),hist=False,kde_kws={'cumulative':True})
sns.distplot(yv_test,hist=False,kde_kws={'cumulative':True})
```

Out[36]: <AxesSubplot:xlabel='y', ylabel='Density'>



```
In [37]: Xv_test['y^'] = modelo_test.predict(Xv_test)
Xv_test['y'] = yv_test
```

```
In [38]: modelo_test.coef_
```

Out[38]: array([-1.48760393e-01, -1.48760393e-01, 1.33688497e-01, -8.21227846e-01,
-1.80719876e-01, 1.03225620e+00, -1.48760393e-01, -1.48760393e-01,

3.47944646e-01,	-1.04221773e-03,	-1.30113183e+00,	-4.62142203e-03,
-1.48760393e-01,	-1.48760393e-01,	2.49133424e-02,	7.01695131e-01,
-7.00091044e-01,	-8.87602136e-02,	-1.24168639e+00,	1.17520706e+00,
1.25832842e-02,	2.84669706e-02,	3.69635076e-01,	-5.15159571e-01,
1.14446785e-11,	-1.80533366e-12,	-7.77923529e-01,	-7.69044832e-01,
2.89467352e+00,	1.30259978e-01,	-1.48760393e-01,	-1.48760393e-01,
-1.48760393e-01,	1.07424143e+00,	5.39737947e-01,	-3.65926858e-01,
-5.71023681e-01,	-2.96726946e+00,	1.29557904e+00,	-1.48760393e-01,
-1.48760393e-01,	-1.48760393e-01,	-4.76540925e-01,	-1.75090748e-01,
2.93042754e-01,	1.09945962e+00,	1.09507451e+00,	-9.74702102e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-7.57862669e-01,
-7.88799788e-01,	-1.52582439e-02,	-1.28742316e+00,	2.97879228e+00,
2.75856969e-01,	-1.70366949e+00,	-1.33331046e+00,	1.45210675e+00,
7.35876946e-02,	3.09158991e-02,	-5.04361851e-02,	1.00563833e-01,
2.19445916e-01,	-1.41264099e-01,	-1.61537450e-14,	-1.33920652e-13,
9.45493683e-14,	1.89428444e+00,	8.11460776e-01,	-5.57794434e-01,
-3.36485940e+00,	-3.24264437e+00,	1.94858723e+00,	-1.48760393e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.22439797e-01,	-9.20638780e-01,
-4.29445645e-01,	-1.84611341e+00,	4.41350368e+00,	1.44743570e+00,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	3.17672330e-02,
1.37878153e-01,	-1.32618603e-01,	2.20882601e-01,	-6.91002117e-01,
-2.05031525e-02,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
1.13389600e-01,	8.67916863e-01,	5.64349997e-01,	2.75574543e+00,
-3.32265811e+00,	-1.58444104e+00,	2.08273892e+00,	2.18622092e+00,
-2.02852314e+00,	1.10005306e-01,	1.05868640e-01,	-1.22938339e-01,
-1.24074186e+00,	-1.23515380e+00,	1.28288192e+00,	-8.54871729e-14,
-1.13242749e-14,	-4.32986980e-14,	-7.15690269e-01,	-7.86635339e-01,
3.68908113e-02,	2.34001202e+00,	3.99790013e+00,	-8.48601770e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
-1.17033693e-01,	-3.50870627e-01,	1.07519623e-01,	5.33245394e-01,
-6.34462692e-01,	1.06644777e+00,	6.06055298e-01,	-1.69160828e+00,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
-2.33149899e-02,	2.14080413e-01,	-4.45045130e-02,	-2.67635320e-02,
4.76397650e-01,	6.29695696e-01,	5.40700948e-01,	-9.24102899e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
3.16897603e-01,	-1.41349375e-02,	-1.90821221e-01,	-2.53143358e-01,
-1.99790477e+00,	-1.08691507e+00,	-1.57045496e+00,	2.78054243e+00,
2.66365062e+00,	2.65850481e+00,	2.66739894e+00,	-2.65055484e+00,
-8.80118936e-02,	-8.24347544e-02,	-1.10679953e-01,	8.67703212e-02,
-4.03816320e-01,	-4.31824601e-01,	-3.14663990e-01,	3.96437487e-01,
-3.17523785e-14,	1.62092562e-14,	-2.98649994e-14,	2.17603713e-14,
-3.79441557e-01,	-7.12657867e-01,	1.79564622e-01,	7.61302476e-01,
2.62448112e-01,	-8.98544562e-01,	9.45667919e-02,	-4.62498233e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
-1.48760393e-01,	-3.69573799e-01,	-4.13698407e-01,	1.13239005e+00,
-2.45156342e-01,	-5.44315796e-01,	9.49376652e-01,	1.69307493e+00,
-4.70068902e+00,	-9.38987518e-01,	2.91975550e+00,	-1.48760393e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
-1.10542426e-01,	2.11426621e-01,	-1.84120913e-01,	-1.62264554e-02,
8.07796667e-02,	7.63924785e-01,	-1.33127957e+00,	-1.41430534e+00,
-1.08195940e+00,	1.43765357e+00,	-1.48760393e-01,	-1.48760393e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	2.53777735e-01,
6.11569094e-02,	-9.08936296e-01,	1.15865762e-01,	4.37881421e-01,
-4.85943772e-01,	6.73858664e-01,	7.44649012e+00,	2.98742216e+00,
-6.81289537e+00,	-1.37684477e+00,	-1.64640277e+00,	-1.56546075e+00,
-1.81046231e+00,	1.44055825e+00,	1.13495202e-01,	1.21187129e-01,
1.35991214e-01,	1.77801935e-01,	-1.17284849e-01,	-9.53549904e-02,
-1.19104619e-01,	-1.64111090e-01,	-3.02896808e-01,	1.11545796e-01,
5.80091530e-15,	-1.33226763e-15,	1.26565425e-14,	-5.32907052e-15,
6.69603262e-15,	-1.71111043e-02,	-4.31129795e-01,	8.54122193e-01,
-2.04420416e-01,	-5.79524192e-01,	-6.89687791e-01,	2.52517627e+00,
-1.33413629e+00,	-9.11990196e-02,	8.68699261e-01,	-1.48760393e-01,
-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,	-1.48760393e-01,
-2.68845444e-01,	-4.30198881e-01,	-8.09855452e-01,	5.15595154e-01,
3.71544952e-01,	2.29458350e+00,	2.46001943e+00,	4.05058079e+00,

```

-1.50003270e+00, -2.84399717e+00, -1.48760393e-01, -1.48760393e-01,
-1.48760393e-01, -1.48760393e-01, -1.48760393e-01, 7.03519969e-03,
-7.25624041e-02, 1.98796634e-01, -3.07641675e-02, -1.08718149e-01,
-3.99784656e-01, -1.23347457e-01, 2.61781404e-01, 6.70140867e-01,
-1.70499938e-01, -1.48760393e-01, -1.48760393e-01, -1.48760393e-01,
-1.48760393e-01, -1.48760393e-01, 5.86854001e-01, 4.22347144e-01,
6.62088394e-01, -1.86258510e-01, -4.58021734e-01, -3.24426286e+00,
-2.56458123e+00, -5.96989635e+00, -2.53422546e-01, 5.30742409e+00,
5.05152213e-01, 5.78096564e-01, 3.89240524e-01, 6.26597780e-01,
-3.50714541e-01, -9.69596369e-02, -9.48210783e-02, -8.78706361e-02,
-1.15754293e-01, 8.14079268e-02, 3.33510099e-01, 3.70735829e-01,
3.19555166e-01, 4.19383341e-01, -3.08197800e-01, -8.88178420e-16,
0.00000000e+00, 3.10862447e-15, -3.77475828e-15, -1.33226763e-15,
-1.99070463e-01, 1.25050175e-01, -5.04444050e-01, 5.06474401e-01,
3.82348449e-02, 2.19918105e+00, 1.19021973e+00, 1.61582891e+00,
-2.16726845e+00, -8.94120871e-01, -1.48760393e-01, 9.33433484e-01,
-2.46596863e+00, -1.48760393e-01, -4.43811363e-01, 5.91878118e-01,
-1.48760393e-01, -1.14077106e+00, 2.18176411e+00, -2.34762466e+00,
1.65298362e-01, -8.15424772e-02, 0.00000000e+00, 1.16371907e+00,
1.35341950e-01, -1.48760393e-01, -1.94399739e-02, 2.39253097e+00,
-1.48760393e-01, 8.96549189e-02, 8.85652794e-01, -1.48760393e-01,
1.38956986e-01, -3.21611261e+00, -1.15765211e+00, 1.69119867e-03,
3.05248060e-01, 0.00000000e+00, -2.71598848e-01, 1.88633632e+00,
-1.48760393e-01, -1.36075873e-01, -3.31376252e+00, -1.48760393e-01,
-6.35789860e-02, 1.34656102e+00, -1.48760393e-01, 2.50974719e-01,
2.69207013e+00, 2.31366120e+00, 1.00786138e-01, -1.25805647e+00,
0.00000000e+00, 1.72419164e-01, -5.60703566e+00, -1.48760393e-01,
7.79090694e-01, -1.00211327e+00, -1.48760393e-01, 1.73270799e-01,
-1.38140334e+00, -1.48760393e-01, -1.15770623e+00, 2.04440461e+00,
2.16136467e+00, -2.87131010e-02, -5.52172746e-01, 0.00000000e+00,
7.53421935e-02, 2.47762726e+00, -1.48760393e-01, 8.44966709e-01,
7.30641236e-01, -1.48760393e-01, 4.05826464e-03, -3.63405248e-02,
-1.48760393e-01, -7.38617825e-01, -1.53793510e+00, -8.27014282e-01,
-3.97980927e-02, 5.59548517e-01, 0.00000000e+00, 7.26728373e-01,
1.23617094e+00, -1.48760393e-01, 8.36611784e-01, -4.70918669e+00,
-1.48760393e-01, -2.49362165e-01, 9.43478171e-01, -1.48760393e-01,
-7.02875995e-01, 5.00561668e+00, -1.57273582e+00, 4.92869140e-02,
1.77723583e-01, 0.00000000e+00, 8.96051179e-01, -5.01175561e+00,
-1.48760393e-01, 2.65063958e-01, -3.83080933e-01, -1.48760393e-01,
5.09553565e-02, 8.35719221e-01, -1.48760393e-01, -2.39717421e-01,
-4.79719337e-01, 2.12053206e+00, 1.27340988e-01, -1.32746825e+00,
0.00000000e+00, -8.11836185e-02, -5.24474048e-02, -1.48760393e-01,
3.24389840e-01, 7.88814318e-02, -1.48760393e-01, 1.65148631e-01,
-4.13702407e-01, -1.48760393e-01, -1.30164145e-01, -1.14939492e+00,
-6.19870349e-01, -6.10749375e-02, 6.09387869e-01, 0.00000000e+00,
2.40205003e-01, 4.95726094e-01])

```

Juntamos todas las funciones para predecir con todos los chunks

```

In [39]: %%time
         number = 1

         for chunk in iterator:

             print(f'Model number {number}')

             chunk = data_filter(chunk)
             chunk = dates(chunk)
             chunk = hours_catalog(chunk)

```

```

hri, hrf = chunk[['id_hour']].describe().T[['min', 'max']].values[0].tolist()
vobs = 24
vdes = 1
anclai = int(hri)+vobs-1
anclaf = int(hrf)-vdes
um = ['PULocationID', 'ancla']
ancla = 24
step = 4
varc = ['trip_distance', 'total_amount', 'n']
vard = ['hora']

print(f'Hra inicial: {hri}, Modelo: {number}')
print(f'Hra final: {hrf}, Modelo: {number}')

X = pd.concat(map(lambda ancla: reduce(lambda x, y: pd.merge(x, y, on=um, how='outer',
    map(lambda k: trans(chunk, ancla, k), range(step, vobs+step, step))), range(
y= pd.concat(map(lambda ancla: target(chunk, ancla, vdes), range(anclai, anclaf+

tad = TAD(X, y, um)
varc = [v for v in tad.columns if v[:2]=="v_"]
tgt = 'y'

miss = 1-tad[varc].describe().T[['count']]/len(tad)

X, Xi = impute(tad, varc)

extreme(X, um, tad)

regression(tad, tgt, Xi)

number += 1

```

Model number 1

Hra inicial: 1.0, Modelo: 1

Hra final: 112.0, Modelo: 1

	variable	ks
3	v_hora_min_total_amount_1_4	0.294406
4	v_hora_min_total_amount_2_4	0.316723
6	v_hora_min_trip_distance_1_4	0.312624
7	v_hora_min_trip_distance_2_4	0.323951
12	v_hora_max_total_amount_1_4	0.312452
..
444	v_hora_sum_n_0_8	0.266582
445	v_hora_sum_total_amount_0_8	0.266818
446	v_hora_sum_trip_distance_0_8	0.266700
448	v_hora_std_total_amount_0_8	0.288464
449	v_hora_std_trip_distance_0_8	0.288464

[156 rows x 2 columns]

```

count    9722.000000
mean      0.358465
std       0.479574
min       0.000000
25%      0.000000
50%      0.000000
75%      1.000000
max       1.000000

```

Name: extremo, dtype: float64

Fitting 10 folds for each of 4 candidates, totalling 40 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.

[Parallel(n_jobs=-1)]: Done 40 out of 40 | elapsed: 3.9s finished

```

14  v_hora_mean_total_amount_0_4  0.345952
..
444          v_hora_sum_n_3_4  0.341100
445  v_hora_sum_total_amount_3_4  0.341568
446  v_hora_sum_trip_distance_3_4  0.341568
448  v_hora_std_total_amount_3_4  0.351450
449  v_hora_std_trip_distance_3_4  0.351450

```

[164 rows x 2 columns]

```

count      9209.000000
mean        0.404061
std         0.490736
min         0.000000
25%         0.000000
50%         0.000000
75%         1.000000
max         1.000000

```

Name: extremo, dtype: float64

Fitting 10 folds for each of 4 candidates, totalling 40 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.

[Parallel(n_jobs=-1)]: Done 40 out of 40 | elapsed: 2.9s finished

18.75571170775452

20.830105257517218

Model number 112

Hra inicial: 1.0, Modelo: 112

Hra final: 30.0, Modelo: 112

	variable	ks
3	v_hora_min_total_amount_2_4	0.195548
6	v_hora_min_trip_distance_2_4	0.225138
12	v_hora_max_total_amount_2_4	0.225138
15	v_hora_max_trip_distance_2_4	0.225138
21	v_hora_mean_total_amount_2_4	0.225138
24	v_hora_mean_trip_distance_2_4	0.225138
27	v_hora_sum_n_2_4	0.225138
30	v_hora_sum_total_amount_2_4	0.225138
33	v_hora_sum_trip_distance_2_4	0.225138
39	v_hora_std_total_amount_2_4	0.249995
42	v_hora_std_trip_distance_2_4	0.249995
49	v_hora_min_total_amount_1_8	0.385702
53	v_hora_min_trip_distance_1_8	0.409316
61	v_hora_max_total_amount_1_8	0.417187
65	v_hora_max_trip_distance_1_8	0.417187
73	v_hora_mean_total_amount_1_8	0.417187
77	v_hora_mean_trip_distance_1_8	0.417187
81	v_hora_sum_n_1_8	0.417187
85	v_hora_sum_total_amount_1_8	0.417187
89	v_hora_sum_trip_distance_1_8	0.417187
97	v_hora_std_total_amount_1_8	0.422749
101	v_hora_std_trip_distance_1_8	0.422749
170	v_hora_min_total_amount_0_16	0.231921
175	v_hora_min_trip_distance_0_16	0.237683
185	v_hora_max_total_amount_0_16	0.237683
190	v_hora_max_trip_distance_0_16	0.239123
200	v_hora_mean_total_amount_0_16	0.239123
205	v_hora_mean_trip_distance_0_16	0.239123
210	v_hora_sum_n_0_16	0.239123
215	v_hora_sum_total_amount_0_16	0.239123
220	v_hora_sum_trip_distance_0_16	0.239123
230	v_hora_std_total_amount_0_16	0.275000
235	v_hora_std_trip_distance_0_16	0.275000
count	640.000000	
mean	0.325000	
std	0.468741	
min	0.000000	
25%	0.000000	

```
50%          0.000000
75%          1.000000
max          1.000000
Name: extremo, dtype: float64
Fitting 10 folds for each of 4 candidates, totalling 40 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 40 out of 40 | elapsed: 0.2s finished
15.399852347320953
38.200627576156414
CPU times: user 7h 51min 52s, sys: 5min 39s, total: 7h 57min 32s
Wall time: 7h 52min 21s
```

Persistencia de los Modelos

```
In [40]: import pickle
```

```
In [45]: with open("models_other.pkl", "wb") as f:
          for model in models:
              pickle.dump(model, f)
```

Si queremos abrir los modelos de nuevo para evitar entrenamiento

```
In [42]: new_models = []
          with open("models.pckl", "rb") as f:
              while True:
                  try:
                      new_models.append(pickle.load(f))
                  except EOFError:
                      break
```

```
In [44]: len(new_models)
```

```
Out[44]: 112
```

```
In [ ]:
```